

# INTRODUCTORY TALK

Dylan Bourgeois ML Engineer Candidate

It's nice to meet  
all of you!



You can find these slides [dtsbourg.me/kodiak](https://dtsbourg.me/kodiak).

# About me



@dtsbourg



contact@dtsbourg.me



dtsbourg.me



## Education

**Student at EPFL (Lausanne, Switzerland)**

B.Sc. in Microengineering

M.Sc. in Robotics

Mainly focused on Machine Learning

## Experience

**Intern at LHCb (CERN)**

Working on designing ML methods for the LHCb Trigger Upgrade. [Bourgeois et al, 2018b, Hasse et al, 2018]

## Latest

**Masters Thesis at Stanford**

Supervised by Prs. J. Leskovec (Stanford, SNAP) & P. Vanderghenst (EPFL, LTS2).

Working on *Learning representations of source code from structure & context*.

MSc. Thesis

---

# Learning representations of source code from structure & context.

*Work done under supervision by*

Pr. Jure Leskoveč / Dr. Michele Catasta (*SNAP, Stanford*)

Pr. Pierre Vanderghenst / Michaël Defferrard (*LTS2, EPFL*)

# Capturing similarities of source code

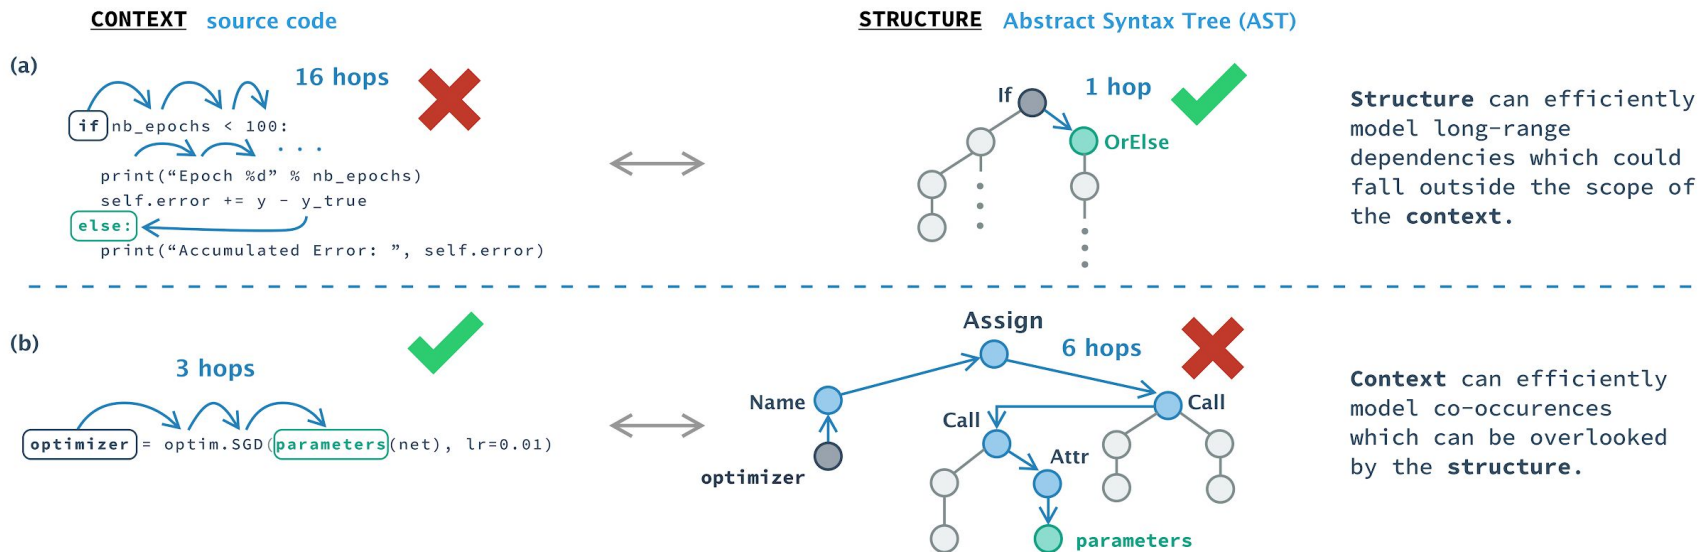
Programming languages offer a unified interface, which is leveraged by programmers. The regularities in coding patterns can be used as a proxy for semantics.

## Example applications

- Code recommendation
- Plagiarism detection
- Smarter development tools
- Error correction
- Smart search

```
for input, target in dataset:  
    optimizer.zero_grad()  
    output = model(input)
```

# Existing approaches only capture a single mode



# We propose a hybrid approach



Our model leverages both **heuristics** and **regularities**, specifically through **structure**.

## HEURISTICS (STRUCTURE)

We provide evidence for the importance of leveraging structure in the representation of source code.

## REGULARITIES (CONTEXT)

We show that patterns in the input provide a decent signal.



## HYBRID (OURS)

We propose a model which learns to recognize both structural and lexical patterns.

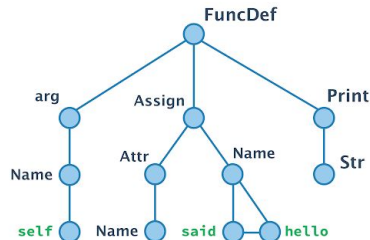
# Structure is readily available for source code

Unlike natural language where parse trees are not unique and costly to infer.

## HEURISTICS (STRUCTURE)

Structure can be extracted deterministically through compiler tools. It represents the language's grammar / syntax.

```
def hello_world(self):  
    self.said_hello = True  
    print("Hello world!")
```



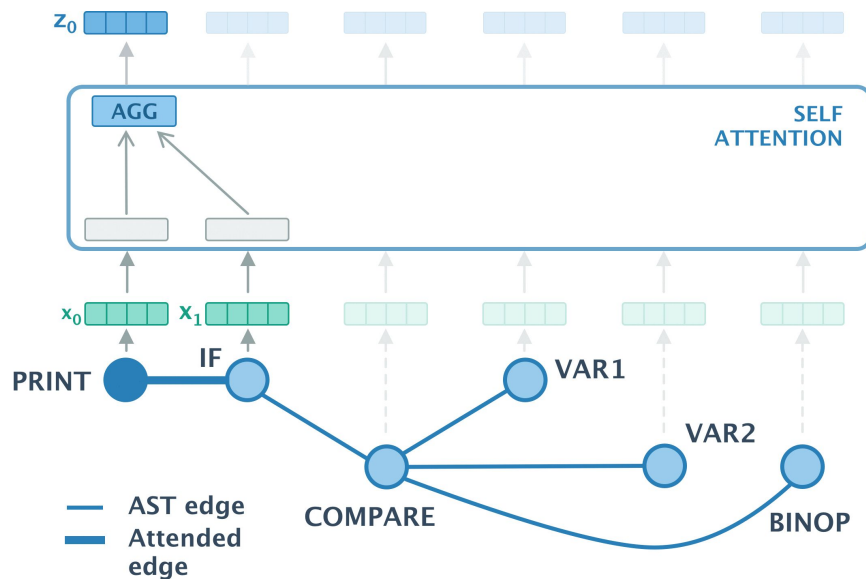
## REGULARITIES (CONTEXT)

We show that patterns in the input provide a decent signal.

## STRUCTURE

# Graph Neural Networks capture local structure

Node representations are computed from each of its neighbours.



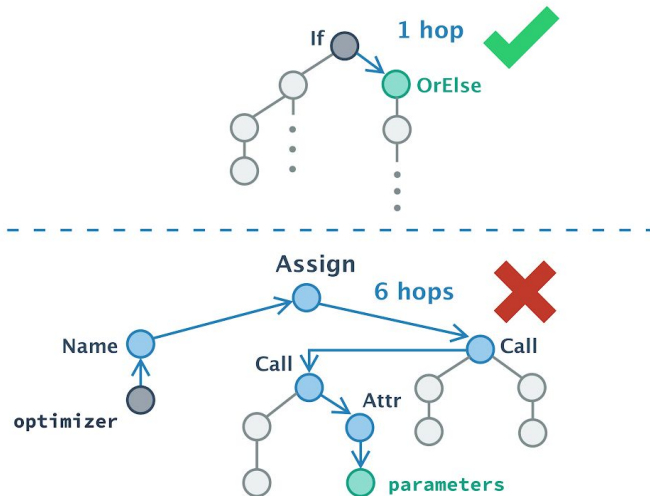


## STRUCTURE

# Graph Neural Networks capture local structure

... but only in a limited receptive field!

### STRUCTURE Abstract Syntax Tree (AST)



# Language Models are learned from context

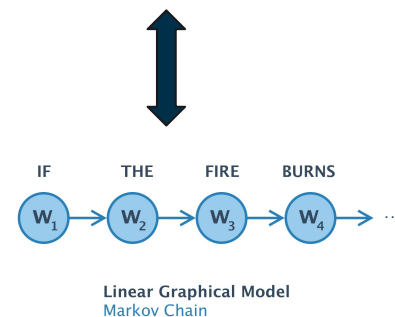
## HEURISTICS (STRUCTURE)

Structure can be extracted deterministically through compiler tools. It represents the language's grammar / syntax.

## REGULARITIES (CONTEXT)

Language models have been tried many times on Natural Language, but also on “Big Code”. [Hindle et al., 2012]

$$P(w_1, \dots, w_m) = \prod_{i=1}^m P(w_i | w_1, \dots, w_{i-1}) \approx \prod_{i=1}^m P(w_i | w_{i-(n-1)}, \dots, w_{i-1})$$

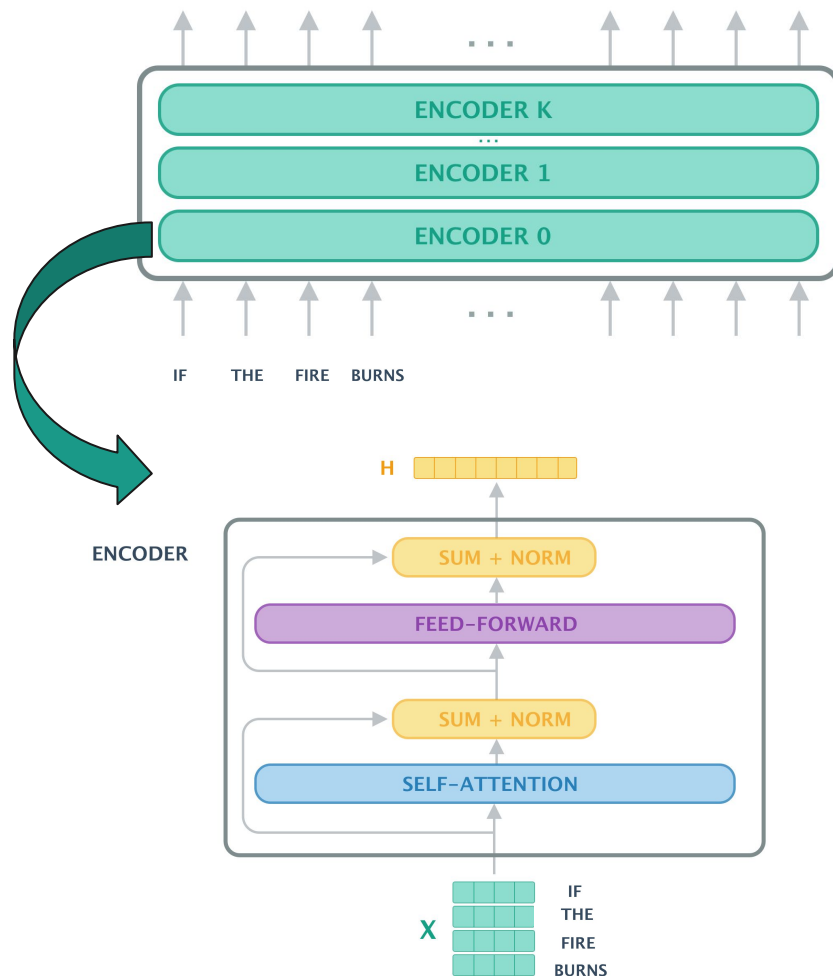


## CONTEXT

# The Transformer is very good at capturing context

No assumptions are made on the underlying structure: the attention module can attend to all the elements in the sequence.

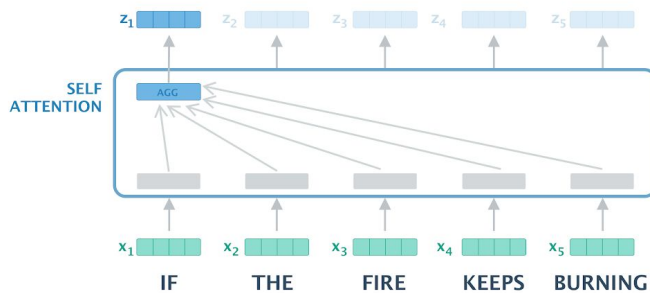
[Vaswani et al., 2017, Devlin et al, 2018]



## INSIGHT

# The Transformer is actually a GNN!

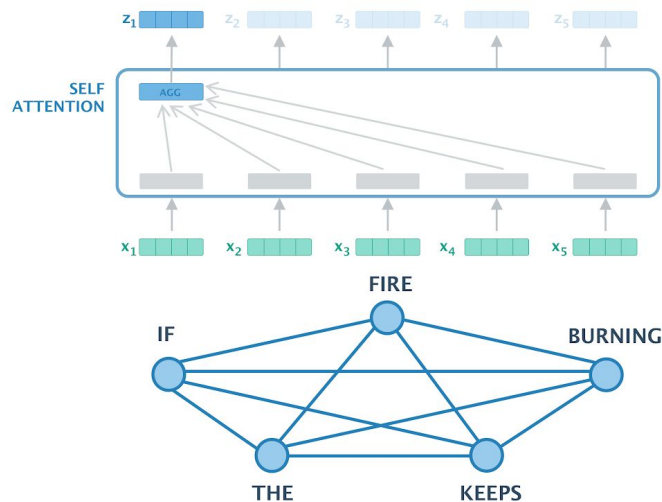
The encoder can be seen as a message-passing Graph Neural Network on a fully connected input graph.



## INSIGHT

# The Transformer is actually a GNN!

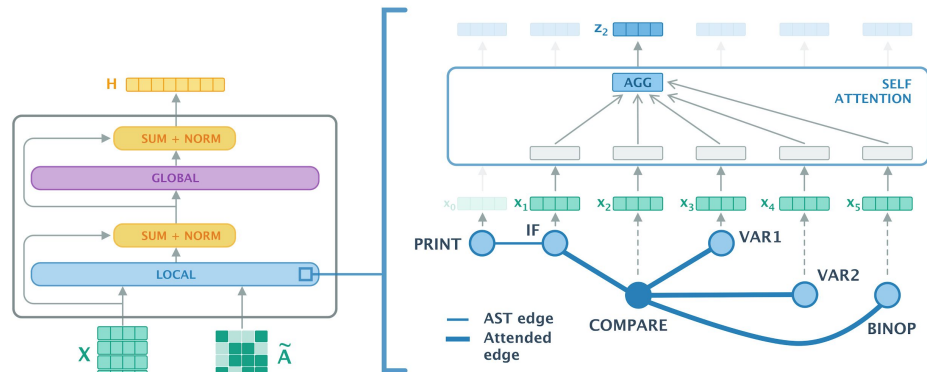
The encoder can be seen as a message-passing Graph Neural Network on a fully connected input graph.



## OUR APPROACH - BiFocale

# Capturing both local structure and global context.

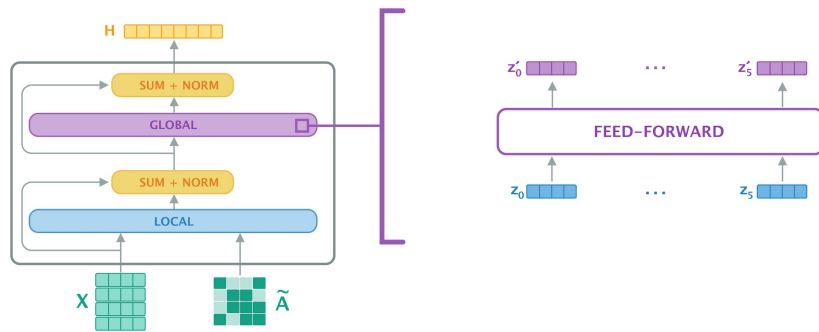
We can modify a Transformer encoder block to run on arbitrarily structured inputs.



## OUR APPROACH - BiFocale

# Capturing both local structure and global context.

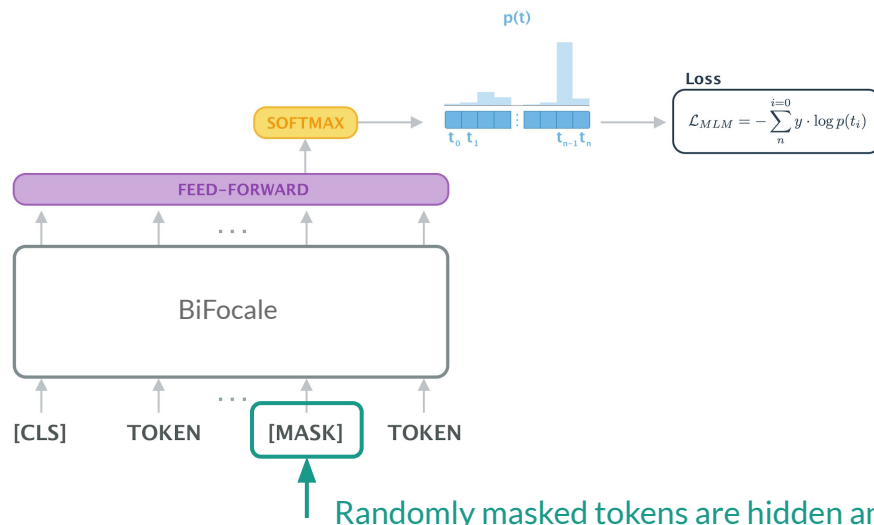
For example, with the masked attention formulation, we can modify a Transformer encoder block to run on arbitrarily structured inputs.



# This hybrid model can also be pre-trained!

[Devlin et al, 2018; Radford et al., 2019]

This has been shown to equip the model with an initial knowledge of the domain at hand. This solution should be closer to future relevant tasks in the hypothesis space.

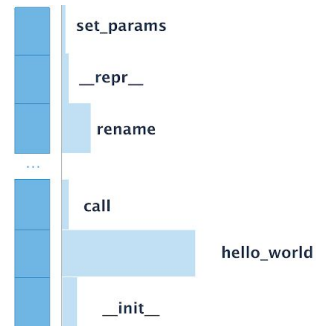
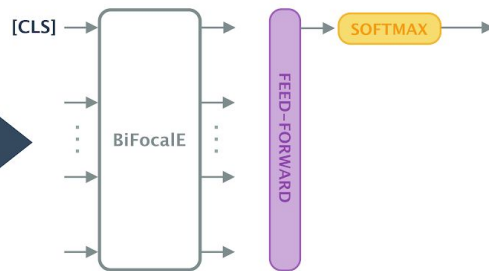
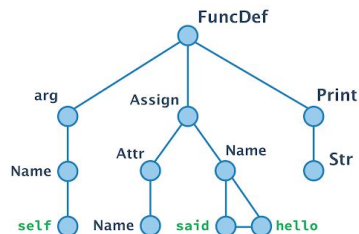






# The model is fine-tuned to predict method names

Given a method definition, the model must predict a relevant name from ~10k.

```
def hello_world(self):  
    self.said_hello = True  
    print("Hello world!")
```



... and is pretty good at it!

	BiFOCALE		Alon' 19 [9]		Alon' 18 [8]	Fernandes' 18 [19]		Allamanis' 18 [7]
	Acc.	F1	Acc.	F1	Acc.	F1	Acc	
JAVA 	<b>0.756</b>	<b>69.1</b>	0.633	59.5	0.473	51.4	—	
PYTHON 	<b>0.760</b>	<b>60.5</b>	—	—	0.511 @7	—	0.416	

Trained on

\* **1M LoC**

\* **3M LoC**

# Predictions hint at the model's hybrid nature

## CORRECT PREDICTIONS

```
1 def get_config(self):  
    return {  
        'mean': self.mean,  
        'stddev': self.stddev,  
        'seed': self.seed  
    }
```

---

**Predictions** 0. get\_config (1.0)  
1. \_updated\_config (0.0)  
2. \_preprocess\_conv3d\_kernel (0.0)

```
2 def __init__(self, minval=0, maxval=5,  
    seed=None):  
    self.minval = minval  
    self.maxval = maxval  
    self.seed = seed
```

---

**Predictions** 0. \_\_init\_\_ (1.0)  
1. on\_train\_begin (0.0)  
2. preprocess\_input (0.0)

## INCORRECT PREDICTIONS

```
1 def glorot_normal(seed=None):  
    return VarianceScaling(scale=1.,  
        mode='fan_avg',  
        distribution='normal',  
        seed=seed)
```

---

**Predictions** 0. he\_normal (0.209)  
1. lecun\_normal (0.198)  
2. lecun\_uniform (0.198)

# BiFocale is a hybrid!

The model can leverage both **co-occurrence based semantics** as well as structural similarities.

```
def sigmoid(x):  
    return 1. / (1. + np.exp(-x))
```

**Predictions** 0. **tanh (0.525)**

```
def tanh(x):  
    return np.tanh(x)
```

1. **softplus (0.335)**

```
def softplus(x):  
    return np.log(1. + np.exp(x))
```

2. **softsign (0.104)**

```
def softsign(x):  
    return x / (1 + np.abs(x))
```

# BiFocale is a hybrid!

The model can leverage both co-occurrence based semantics as well as **structural similarities**.

```
def sigmoid(x):  
    return 1. / (1. + np.exp(-x))
```

**Predictions** 0. **tanh (0.525)**

```
def tanh(x):  
    return np.tanh(x)
```

1. **softplus (0.335)**

```
def softplus(x):  
    return np.log(1. + np.exp(x))
```

2. **softsign (0.104)**

```
def softsign(x):  
    return x / (1 + np.abs(x))
```





## BiFocale sets a new SoTA!

We also show how heuristics are useful for learning on structured grammars.

	Accuracy			
	@1	@3	@5	@7
BERT	0.3	0.43	0.48	0.52
BiFOCALE	<b>0.59</b>	<b>0.792</b>	<b>0.833</b>	<b>0.849</b>
Alon et al.'18 [8]	—	—	—	0.567
Allamanis et al.'18 [7]	0.323	0.408	0.437	—

# Predictions hint at the model's hybrid nature

## CORRECT PREDICTIONS

```
1 for cell in self.cells:
    if isinstance(cell, Layer):
        trainable_weights += cell.trainable_weights
```

### Predictions

```
['cell', '[PAD]', '[PAD]', '[PAD]']
```

```
2 def __call__(self, shape, dtype=None):
    return K.constant(0, shape=shape,
                      dtype=dtype)
```

### Predictions

```
['self', '[PAD]', '[PAD]', '[PAD]']
```

## INCORRECT PREDICTIONS




```
1 def top_k_categorical_accuracy(y_true,
                                y_pred, k=5):
    return K.mean(K.in_top_k(y_pred,
                              K.argmax(y_true, axis=-1), k),
                  axis=-1)
```

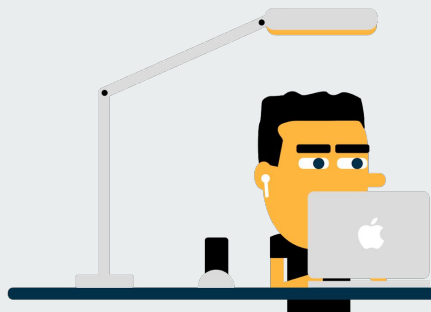
### Predictions

```
0. ['y', 'true', 'true', 'true']
1. ['self', '[PAD]', '[PAD]', '[PAD]']
2. ['true', 'train', 'train', 'train']
```



# For more ...

 @dtsbourg  
 contact@dtsbourg.me  
 dtsbourg.me



Stanford

## GNN Explainer

Interpretability methods for Graph Neural Networks. [Ying et al, 2019]

CERN

## LHCb Trigger Upgrade (CERN)

Working on designing ML methods for the LHCb Trigger Upgrade. [Bourgeois et al, 2018b, Hasse et al, 2018]

EPFL

## Media Observatory

Monitoring the media ecosystem. [Bourgeois et al, 2018a, Rappaz et al, 2019]

CERN

## **LHCb Trigger Upgrade (CERN)**

Working on designing ML methods for the LHCb Trigger Upgrade. [Bourgeois et al, 2018b, Hasse et al, 2018]

EPFL

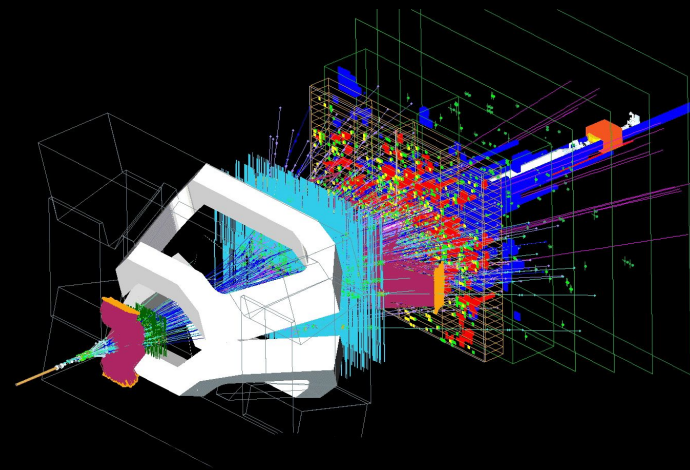
## **Media Observatory**

Monitoring the media ecosystem. [Bourgeois et al, 2018a, Rappaz et al, 2019]

[Bourgeois et al, 2018b, Hasse et al, 2018]

## Fast selection of interesting collisions

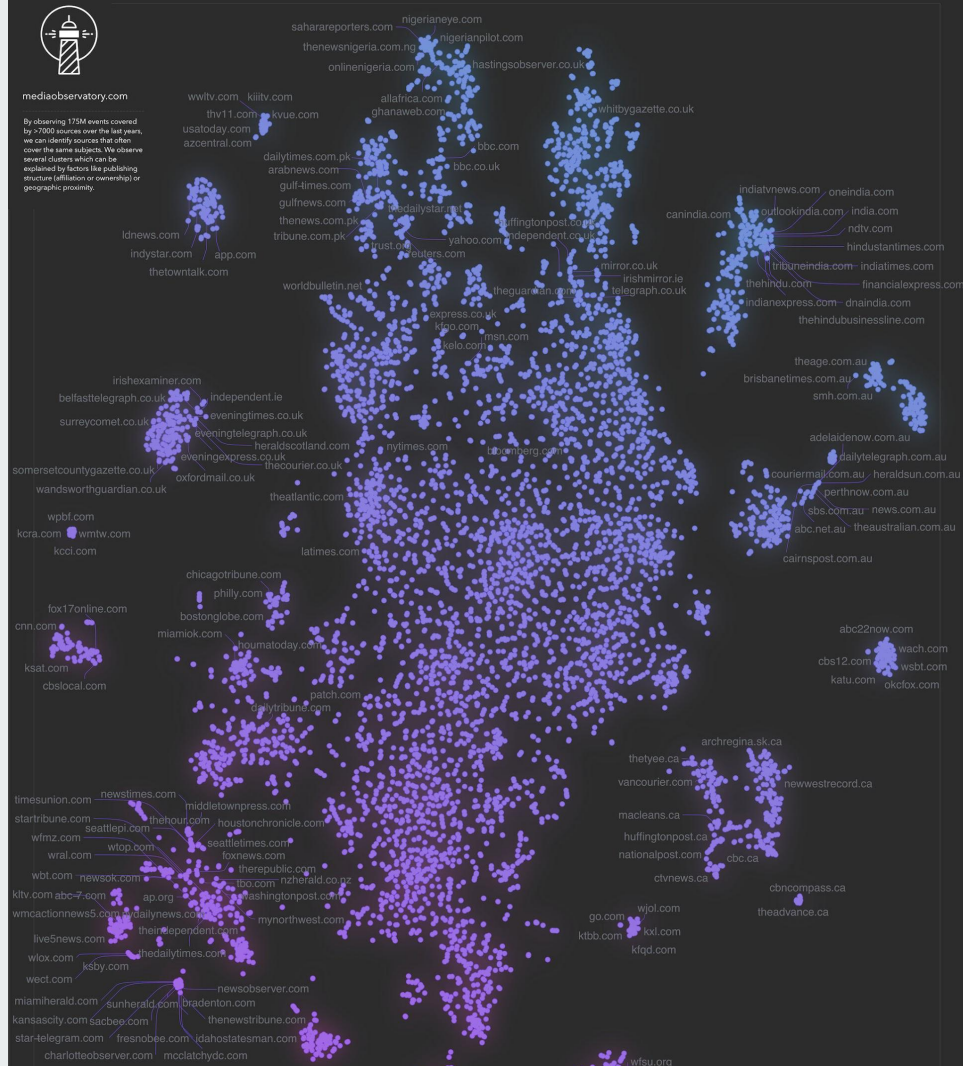
The experiment would be throttled if all events were saved to disk so the Trigger acts as a filter.



[Bourgeois et al, 2018a, Rappaz et al, 2019]

# Studying coverage patterns to observe the news ecosystem

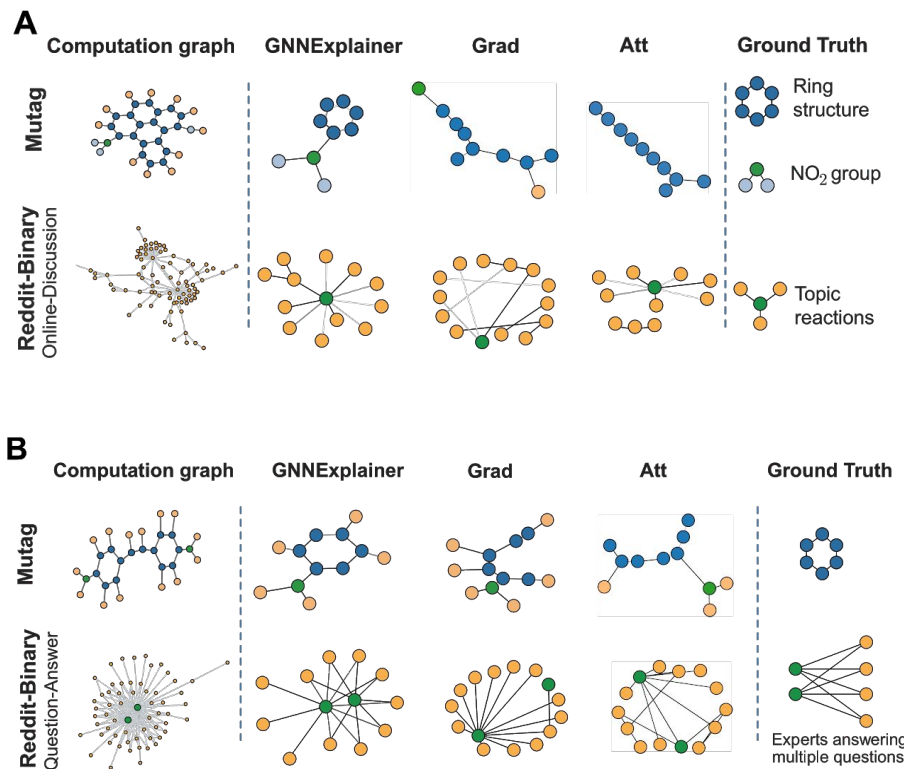
We build a map of source similarity without relying on a ground truth to estimate bias.



[Ying et al, 2019]

# Interpretability methods for Graph Neural Networks

Works for any trained GNN, showing important structural and feature-based information that is most relevant to a prediction.



—

# Thank you ~~K~~odiak!

Questions?

**Dylan Bourgeois**  
@dtsbourg



# Bibliography – My work

---

**[Ying et al, 2019]** *GNN Explainer: A Tool for Post-hoc Explanation of Graph Neural Networks* \*

R. Ying, D. Bourgeois, J. You, M. Zitnik, J. Leskovec

**[Bourgeois et al, 2018a]** *Selection Bias in News Coverage: Learning It, Fighting It*

D. Bourgeois, J. Rappaz, K. Aberer, WWW'18

**[Rappaz et al, 2019]** *A Dynamic Embedding Model of the Media Landscape*

J. Rappaz, D. Bourgeois, K. Aberer, WWW'19

**[Bourgeois et al, 2018b]** *Using holistic information in the Trigger*

D. Bourgeois, C. Fitzpatrick, S. Stahl, LHCb Pub

**[Hasse et al, 2018]** *New approaches for track reconstruction in LHCb's Vertex Locator*

C. Hasse, J. Albrecht, B. Couturier, D. Bourgeois, V. Coco, N. Nolte, S. Ponce, JHEP'18

\*

Submitted to NeurIPS under the title *GNN Explainer: Generating Explanations for Graph Neural Networks*.

# Bibliography

---

**[Allamanis, 2018]** Allamanis, M. (2018). *The adverse effects of code duplication in machine learning models of code*. arxiv:1812.06469.

**[Allamanis et al., 2015]** Allamanis, M., Barr, E. T., Bird, C., and Sutton, C. (2015). *Suggesting accurate method and class names*. ESEC/FSE 2015, pages 38–49

**[Allamanis et al., 2018a]** Allamanis, M., Barr, E. T., Devanbu, P. T., and Sutton, C. A. (2018a). *A survey of machine learning for big code and naturalness*. ACM Comput. Surv., 51:81:1–81:37.

**[Allamanis et al., 2018b]** Allamanis, M., Brockschmidt, M., and Khademi, M. (2018b). *Learning to represent programs with graphs*. ICLR.

**[Alon et al., 2018]** U. Alon, M. Zilberstein, O. Levy, and E. Yahav. *A general path-based representation for predicting program properties*. PLDI 2018.

**[Alon et al., 2019]** Alon, U., Zilberstein, M., Levy, O., and Yahav, E. (2019). *Code2vec: Learning distributed representations of code*. POPL.



# Bibliography

---

**[Bengio et al., 2003]** Bengio, Y., Ducharme, R., Vincent, P., and Janvin, C. (2003). *A neural probabilistic language model*. J. Mach. Learn. Res., 3:1137–1155.

**[Collobert and Weston, 2008]** Collobert, R. and Weston, J. (2008). *A unified architecture for natural language processing: Deep neural networks with multitask learning*. ICML '08.

**[Deerwester et al., 1990]** Deerwester, S.C., Dumais, S.T., Landauer, T.K., Furnas, G.W., and Harshman, R. A. (1990). *Indexing by latent semantic analysis*. JASIS, 41:391–407.

**[Devlin et al., 2018]** Devlin, J., Chang, M.-W., Lee, K., and Toutanova, K. (2018). *BERT: Pre-training of deep bidirectional transformers for language understanding*. Arxiv:1810.04805.

**[Firth, 1957]** Firth, J.R. (1957). *A synopsis of linguistic theory 1930-55*. Studies in Linguistic Analysis (special volume of the Philological Society), 1952-59:1–32.

**[Fernandes, 2018]** P. Fernandes, M. Allamanis, and M. Brockschmidt. *Structured neural summarization*, 2018.

**[Hindle et al., 2012]** Hindle, A., Barr, E. T., Su, Z., Gabel, M., and Devanbu, P. (2012). *On the naturalness of software*. In ICSE '12, pages 837–847, IEEE Press.

# Bibliography

---

**[Hamilton et al., 2017]** William L. Hamilton, Rex Ying, Jure Leskovec. Inductive Representation Learning on Large Graphs. NeurIPS 2017

**[Mikolov et al., 2013]** Mikolov, T., Chen, K., Corrado, G., and Dean, J. (2013). *Efficient estimation of word representations in vector space*. ICLR'13

**[Radford et al., 2018]** Radford, A., Narasimhan, K., Salimans, T., and Sutskever, I. (2018). *Improving language understanding by generative pre-training*. OpenAI.

**[Shannon, 1950]** Shannon, C. (1950). *Prediction and entropy of printed english*. Bell Systems Technical Journal.

**[Vaswani et al., 2017]** Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, L., and Polosukhin, I. (2017). *Attention is all you need*. In NeurIPS.

**[Xu et al., 2019]** Xu, K., Hu, W., Leskovec, J., and Jegelka, S. (2019). *How powerful are graph neural networks?* In ICLR'19.